

2024



Elastic Helps Users Improve Security Posture and Productivity Using Amazon Bedrock and Anthropic’s Claude 3 Models

Learn how Elastic uses Amazon Bedrock to help customers improve security analyst efficiency.

[Overview](#) | [Opportunity](#) | [Solution](#) | [Outcome](#) | [AWS Services Used](#) | [Architecture Diagram](#)



Hi, I can connect you with an AWS representative or answer questions you have on AWS.



Reduced

average time to respond

Saved

in general



Need more info? Highlight any text to get an explanation generated with AWS generative AI.



Reduced hallucinations

Provides access



Overview

[Elastic](#) aims to help customers find information faster while running their applications smoothly and protecting against cyberthreats. Operating from the premise that search functionality is fundamental to critical business processes, Elastic offers solutions in search, observability, and security, built on Amazon Web Services (AWS).

Recognizing the power of generative artificial intelligence (AI) to better understand user queries, Elastic offers a choice of foundation models (FMs) and large language models (LLMs) to deliver the most relevant results. For instance, Elastic's users can automatically access Anthropic's cutting-edge family of LLMs, including the [Claude 3](#) family of models. To provide this model choice to its users, Elastic uses [Amazon Bedrock](#), a fully managed service that offers users a choice of high-performing FMs from leading AI companies, like Anthropic. Amazon Bedrock also provides a broad set of capabilities to build [generative AI applications](#) with security, privacy, and responsible AI. Using Anthropic's Claude 3 models in Amazon Bedrock, Elastic's customers save hours identifying, remediating, and reporting on security issues, ultimately improving their overall security posture and productivity.



Opportunity | Using Anthropic's Claude 3 Models in Amazon Bedrock to Increase Efficiency for Security Analysts

Elastic, an [AWS Partner](#), supports conventional keyword and text-based searches as well as vector search capabilities using its AI-powered [Elasticsearch Relevance Engine](#) (ESRE), which runs on AWS. "The number-one challenge security teams face today is figuring out the meaning behind all the alerts that are coming in," says Spiteri, Elastic's director of product management for security. Powered by ESRE, the security a



[Elastic Security](#) aims to help users make sense of security alerts and alleviate time-consuming and complex security operations tasks.

Elastic wanted to provide a user-friendly way to complement its threat detection analytics, cloud security, and endpoint protection capabilities. Elastic customers need the ability to select which high-performing model is best for a particular use case while maintaining best practices in privacy and security. Additionally, through its patent-pending [Attack Discovery](#) feature, Elastic provides users the ability to automatically sort hundreds of alerts to identify which are significant attacks.

Solution | Increasing Productivity and Saving Hours for Security Analysts Using Amazon Bedrock

Using generative AI, the [Elastic AI Assistant](#) helps security operations teams with alert triage and investigation, provides suggestions for optimal code efficiency, and converts queries from other platforms into Elastic syntax. ESRE uses retrieval-augmented generation (RAG) to fetch company data to enrich prompts and responses while considering the privacy and permissions of the underlying content. In Amazon Bedrock, users can ask Anthropic's Claude questions in natural language, such as: "Can you summarize what has happened in my environment today?" By talking to large bodies of content, Claude can summarize, perform question and answering, forecast trends, compare multiple documents, and much more.

From a simple drop-down menu, users can automatically access the latest model of Anthropic's Claude and relay a large volume of information to the model, which provides a 200,000-token context window, equivalent to around 150,000 words or 500 pages of text. "We benefit from the scalability, performance, accuracy, and privacy of Amazon Bedrock," says Uday Theepireddy, senior principal solutions architect at Elastic. "Because Amazon Bedrock is a fully managed service from AWS, we don't worry about updating or customizing models. We can rely on Amazon Bedrock to have the latest and most cutting-edge FMs in the world available to us and our customers."

Using [Model Evaluation on Amazon Bedrock](#), Elastic customers can compare high-performing LLMs using automatic evaluation with predefined metrics such as accuracy, robustness, and toxicity or human evaluation through subjective metrics such as friendliness or style. "The ability to select from the top-performing FMs like Anthropic's Claude 3 models in Amazon Bedrock has been extremely beneficial," says Spiteri. "We can compare the outputs, metrics, and associated costs across different models so that we make an informed decision on which model would be most suitable for what we are trying to accomplish. This has significantly streamlined our process and saved a considerable time in deploying our applications."

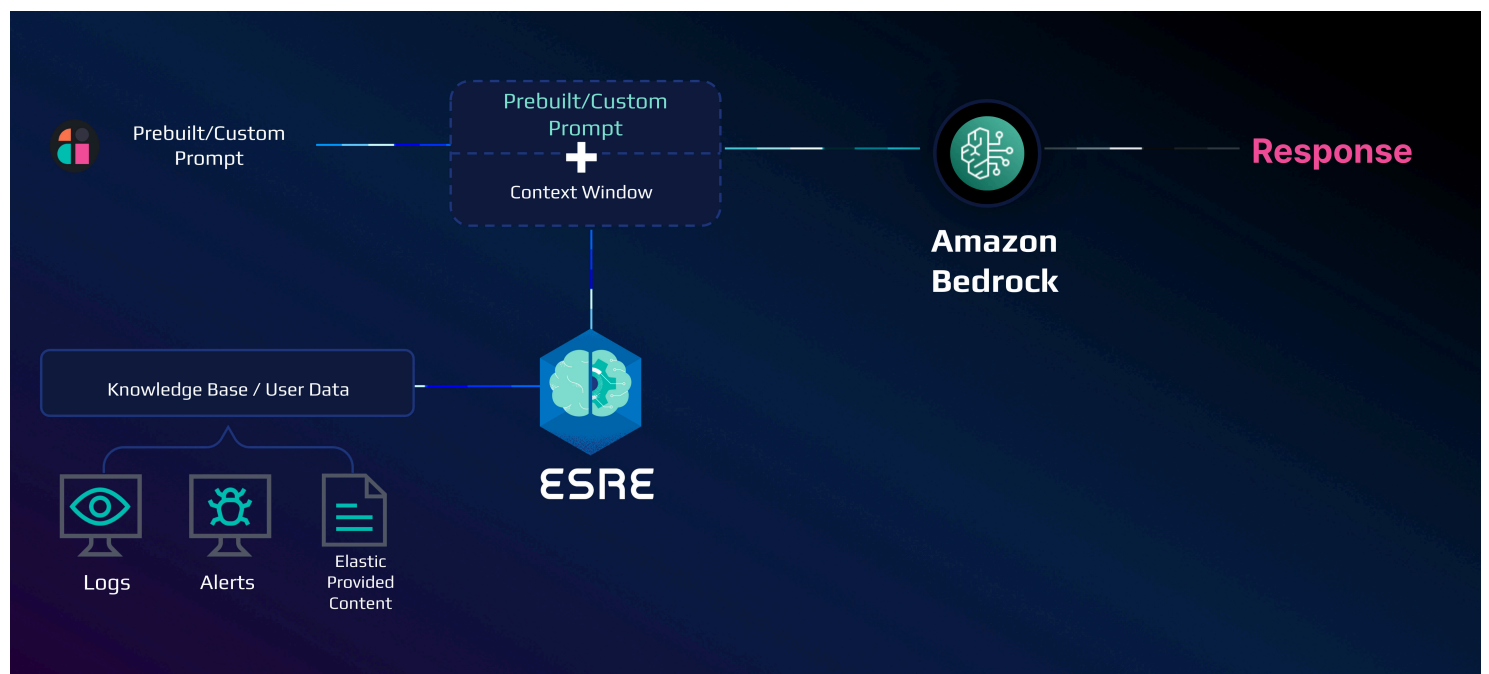


Elastic customers using Amazon Bedrock default to [AWS PrivateLink](#), which provides private connectivity between virtual private clouds, supported AWS services, and on-premises networks without exposing customers' traffic to the public internet. "We wouldn't have integrated Amazon Bedrock if we didn't have confidence in its privacy policy and security permissions," says Spiteri. Customers using Amazon Bedrock through Elastic adjust to data residency requirements by establishing connectors from multiple [AWS Regions](#). "Using Amazon Bedrock, we are changing the game for security investigations," says Spiteri. "It's no longer cat and mouse; our customers get answers more quickly and in context of why they asked."

Using Amazon Bedrock to access Anthropic's latest state-of-the-art Claude models, the Elastic AI Assistant can review security logs, analyze data, and accurately answer questions in seconds. "Having access to Claude in Amazon Bedrock with the Elastic AI Assistant helps teams close the blind spots and give details about the security landscape that they might not see as a human," says Theepireddy. As a result, customers significantly reduce the average time to respond, saving analysts hours. Analysts can now write reports and generate queries more quickly, especially for users who are unfamiliar with Elastic's syntax. "I can't talk enough about how helpful the Elastic AI Assistant, built on Amazon Bedrock and Anthropic's Claude, has been for our users," says Spiteri. "We are always discovering new ways of streamlining, adding more functionality, and solving more problems using Amazon Bedrock."

Elastic's Attack Discovery feature also uses Claude 3. Attack Discovery analyzes patterns and correlates events in near real time, displaying pertinent information for users in an intuitive interface. In addition to productivity gains, security teams have improved incident management.

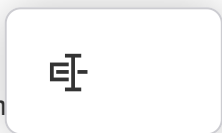
Architecture Diagram



Outcome | Accessing High-Performing Models in Amazon Bedrock to Lower Risk

In the future, Elastic will provide its customers access to even more FMs through Amazon Bedrock. For example, Elastic plans to use Amazon Bedrock to access additional high-performing models that support image recognition, voice-to-text conversion, and multilingual capabilities.

"Using AWS, we give our customers the right tools to get the job done," says Spiteri. "Powering through LLMs such as the Claude 3 family of models through Amazon Bedrock definitely helps us solve our main



mission: to help our customers lower their risk in the most cost-effective and efficient way possible.”

About Elastic

Founded in 2012, Elastic (formerly Elasticsearch) is an American-Dutch company that offers software solutions in enterprise search, observability, and security.

AWS Services Used

Amazon Bedrock

Amazon Bedrock is a fully managed service that offers a choice of high-performing foundation models (FMs) from leading AI companies like AI21 Labs, Anthropic, Cohere, Meta, Mistral AI, Stability AI, and Amazon through a single API, along with a broad set of capabilities you need to build generative AI applications with security, privacy, and responsible AI.

[Learn more »](#)

AWS PrivateLink

AWS PrivateLink provides private connectivity between virtual private clouds (VPCs), supported AWS services, and your on-premises networks without exposing your traffic to the public internet. Interface VPC endpoints, powered by PrivateLink, connect you to services hosted by AWS Partners and supported solutions available in AWS Marketplace.

[Learn more »](#)

More Generative AI Customer Stories

ITALY

BELGIUM





Namirial Streamlines Contract Signing with...

Namirial enhanced its eSianAnvwhere



2024



ElysianNxt Uses Generative AI and AWS to Help...

ElvsianNxt built a risk management



2024

COSTA RICA



Modernizing Applications in Minutes Using Amazon Q...

Learn how Novacomp accelerated



2024

UNITED STATES



Prime Video's new feature uses generative AI to...

X-Ray Recaps conveniently catches you up on what's happening, whenever you're in need of a refresher.

2024

1 2 3 4 5 6 ... 55

Get Started

Organizations of all sizes across all industries are transforming their businesses and delivering on their missions every day using AWS. Contact our experts and start your own AWS journey today.

Contact Sales



2

Learn About AWS

- What Is AWS?
- What Is Cloud Computing?
- AWS Accessibility
- AWS Inclusion, Diversity & Equity
- What Is DevOps?
- What Is a Container?
- What Is a Data Lake?
- What is Artificial Intelligence (AI)?
- What is Generative AI?
- What is Machine Learning (ML)?
- AWS Cloud Security
- What's New
- Blogs
- Press Releases

Resources for AWS

- Getting Started
- Training and Certification
- AWS Solutions Library
- Architecture Center
- Product and Technical FAQs
- Analyst Reports
- AWS Partners

Developers on AWS

- Developer Center
- SDKs & Tools
- .NET on AWS
- Python on AWS
- Java on AWS
- PHP on AWS
- JavaScript on AWS

Help

- Contact Us
- Get Expert Help
- File a Support Ticket
- AWS re:Post
- Knowledge Center
- AWS Support Overview
- Legal
- AWS Careers



Amazon is an Equal Opportunity Employer: *Minority / Women / Disability / Veteran / Gender Identity / Sexual Orientation / Age.*

Language

- عربي |
- Bahasa Indonesia |
- Deutsch |
- English |
- Español |
- Français |
- Italiano |
- Português |
- Tiếng Việt |
- Türkçe |
- Русский |
- ไทย |
- 日本語 |
- 한국어 |

