



Healthify Innovates Using AWS Generative AI Solutions and Achieves Operational Excellence with AWS Enterprise On-Ramp

Learn how healthtech company Healthify developed AI solutions more efficiently using Amazon Bedrock and achieved operational excellence using AWS Enterprise On-Ramp.

[Overview](#) | [Opportunity](#) | [Solution](#) | [Outcome](#) | [AWS Services Used](#)

40%

reduction in costs

30%

reduction in query response times | uptime maintained

Hours to minutes

reduced developer time



Hi, I can connect you with an AWS representative or answer questions you have on AWS.



Need more info? Highlight any text to get an explanation generated with AWS generative AI.



Overview

Health and wellness service provider [Healthify](#) has been using Amazon Web Services (AWS) since the company's inception in 2012, and it had been building its own artificial intelligence (AI) solutions since 2017. As its operations grew, the company adopted additional AWS solutions as part of an initiative to make developing new AI features more efficient, increase profitability, and achieve operational excellence.

Using AWS, Healthify reduced costs and query times, maintained uptime, decreased developer time from hours to minutes, and accelerated AI model deployment.



Opportunity | Developing Generative AI Solutions Using AWS

Healthify is one of the largest healthtech companies globally, with about 40 million registered users on its wellness application. Its goal is to improve the health of 1 billion people worldwide by focusing on the four pillars of nutrition, fitness, stress, and sleep. Healthify's mobile app has lifestyle trackers, access to nutritionist coaches and fitness trainers, and personalized plans to help users achieve their health goals.

Healthify turned to AWS to create an AI assistant for its app and give users access to AI-generated fitness plans. The AI and machine learning (ML) collaboration began in 2018 when Healthify began using [Amazon SageMaker](#)—a service to build, train, and deploy ML models for any use case—to build its AI/ML applications with faster iteration times. In 2022, Healthify integrated SageMaker capabilities in its MLOps flow—by migrating the MLOps to Amazon Sagemaker, Healthify can deploy and test a new model in 3–4 hours compared with 2–3 days before. Also, Healthify reduced the time needed to adapt the model by ~50 percent. With the flexibility to automate workflows, Healthify reduced its iteration times from hours to minutes.

In 2023, Healthify began using [Amazon Bedrock](#), a fully managed service that offers a choice of high-performing foundation models from leading AI companies. It deployed its conversation intelligence solution, which extracted key insights from audio calls in multiple Indian languages through autonomous emotion and sentiment analysis. With an average of over 1,600 calls analyzed every month, the call transcripts summaries help coaches provide superior customer service. Additionally, auditors could use the in-depth analysis to identify trends, themes, and potential training opportunities.



The conversation intelligence solution uses Amazon SageMaker for speaker diarization, transcribing, and translating audio files using open-source ML models. It employs [Amazon Comprehend](#), a natural language processing service, to extract sentiments and specific entities from the transcripts. Furthermore, the solution uses generative AI models, such as Anthropic Claude on Amazon Bedrock, to summarize conversations, identify action items, detect issues, and monitor key performance indicators.

“Our Amazon Bedrock–powered Call Analytics solution is expected to provide up to a 75 percent increase in our nutritionist coaches’ capacity to accommodate clients,” says Abhijit Khasnis, vice president of engineering at Healthify.

Healthify sought ways to integrate richer solutions for better customer service while increasing profitability and optimizing operations. One example is Healthify’s rollout of Snap in early 2023, a photo-based food recognition system that makes nutrition tracking easier by using vision-based ML models to detect food items. Snap used open-source vision models and trained and deployed them using Amazon SageMaker. The model can detect over 10,000 varieties of Indian food in single food photos, and the team continues to tune the model’s accuracy.

Solution | Optimizing Infrastructure and Reducing AWS Spend by 40% Using AWS Enterprise On-Ramp

To support business growth and provide an uninterrupted experience for users, Healthify subscribed to [AWS Enterprise On-Ramp](#)—where businesses get technical guidance from a pool of Technical Account Managers (TAMs), service experts, and cloud engineers. Healthify used AWS Enterprise On-Ramp for several optimization initiatives, such as increasing its infrastructure’s ability to handle high traffic volumes.

For instance, every January, the company experiences a traffic increase of nearly two and a half times the traffic of November and December combined. Healthify worked with its TAM to review critical resources ahead of peak periods, proactively identifying and mitigating potential risks. Previously, Healthify had to overprovision infrastructure to meet scaling needs, incurring heavy costs. Using AWS Enterprise On-Ramp, Healthify improved resilience and scaled to support 50 percent more traffic while limiting cost increase to 10 percent.

“In January 2024, our systems that run on AWS had 99.99 percent uptime,” says Manjunath DS, lead engineer at Healthify. “From the perspective of the infrastructure and data teams, we have peace of mind that we will be able to resolve issues quickly.”



Using AWS Enterprise On-Ramp, Healthify optimized its AWS spend as it expanded across regions and onboarded new users. The TAM team identified opportunities to reduce underused resources, migrate to GP3, transition from third-party to AWS tools for managing Spot Instances, and optimize data transfer charges. Healthify implemented these recommendations to reduce AWS spend by approximately 40 percent.

Healthify had been using a third-party extract, load, transform tool but experienced issues due to its dynamic requirements. So, it migrated to [AWS Database Migration Service](#) (AWS DMS), a managed migration and replication service that helps move your database and analytics workloads to AWS. Also, Healthify improved the efficiency of using Amazon Redshift—a cloud data warehouse that businesses use to power data-driven decision —after using AWS Enterprise On-Ramp to perform a deep-dive review of its cluster. Overall, it achieved 30 percent reduction in query response times, reducing the operational bandwidth to manage the cluster.

“When we review our internal metrics, we see that all our analyst queries have far lower waiting times,” says Anil Nayak, lead engineer at Healthify. “We were able to explore [Amazon Redshift](#) features that we use to distribute queries properly and add compute whenever needed, which helps our team work more efficiently.”

Healthify also used [AWS Countdown](#), a service designed for a broad range of cloud use cases, including migrations, modernizations, product launches, streaming, and go-live events, offered under AWS Enterprise On-Ramp to upgrade critical instances such as RDS MySQL 8.0 and Redshift RA3. The upgrade was performed in phases with minimal business impact. The AWS Enterprise On-Ramp TAM team prepared Healthify engineers for common issues seen in similar upgrade scenarios. The TAM team was on standby during live upgrades to swiftly address any issues encountered, offering critical technical support in peak traffic times.

Outcome | Expanding around the World

Healthify aims to continue expanding internationally and is looking at AWS solutions that can achieve high availability and low latency in all geographic regions. It plans to keep innovating in the generative AI space and partnering with the TAM team as it scales to onboard more users, improve operational maturity, promote organizational cost awareness, and strengthen security posture.

“By using AWS, we can stay at the forefront of evolving technology to deliver innovative health and wellness services at scale,” says Khasnis. “Over the last decade, AWS has been a cornerstone of our cloud infrastructure as we expanded our services. We’ve found immense value in the ML support provided by AWS in helping us craft generative AI sample implementations. Furthermore, AWS Enterprise Support On-Ramp has proven invaluable, offering critical technical support during peak traffic times.”

About Healthify

Healthify is an Indian healthcare and wellness service provider with over 340 million customers globally. Its mission is to make healthy living simple, accessible, and affordable for everyone through its nutrition and fitness services.

AWS Services Used



AWS Enterprise On-Ramp

With Enterprise On-Ramp, you get access to many of the features of Enterprise Support, including 24x7 technical support from elite engineers, fast response times, and tools and technology to automatically manage the health of your environment.

[Learn more »](#)

Amazon SageMaker

Amazon SageMaker is built on Amazon's two decades of experience developing real-world ML applications, including product recommendations, personalization, intelligent shopping, robotics, and voice-assisted devices.

[Learn more »](#)

AWS Countdown

Optimize your business-critical events, product launches, migrations, and modernizations on AWS.

[Learn more »](#)

Amazon Bedrock

A fully managed service that offers a choice of high-performing foundation models (FMs) from leading AI companies like AI21 Labs, Anthropic, Cohere, Meta, Mistral AI, Stability AI, and Amazon through a single API, along with a broad set of capabilities you need to build generative AI applications with security, privacy, and responsible AI.

[Learn more »](#)

Get Started

Organizations of all sizes across all industries are transforming their businesses and delivering on their missions every day using AWS. Contact our experts and start your own AWS journey today.

Contact Sales

Learn About AWS

- What Is AWS?
- What Is Cloud Computing?
- AWS Accessibility
- AWS Inclusion, Diversity & Equity
- What Is DevOps?
- What Is a Container?
- What Is a Data Lake?
- What is Artificial Intelligence (AI)?
- What is Generative AI?
- What is Machine Learning (ML)?
- AWS Cloud Security
- What's New
- Blogs
- Press Releases

Resources for AWS

- Getting Started
- Training and Certification
- AWS Solutions Library
- Architecture Center
- Product and Technical FAQs
- Analyst Reports
- AWS Partners

Developers on AWS

- Developer Center
- SDKs & Tools
- .NET on AWS
- Python on AWS
- Java on AWS
- PHP on AWS
- JavaScript on AWS

Help

- Contact Us
- Get Expert Help
- File a Support Ticket
- AWS re:Post
- Knowledge Center
- AWS Support Overview
- Legal
- AWS Careers



Amazon is an Equal Opportunity Employer: *Minority / Women / Disability / Veteran / Gender Identity / Sexual Orientation / Age.*

Language

عربي |



Bahasa Indonesia |
Deutsch |
English |
Español |
Français |
Italiano |
Português |
Tiếng Việt |
Türkçe |
Русский |
ไทย |
日本語 |
한국어 |
中文 (简体) |
中文 (繁體)

Privacy
|
Accessibility
|
Site Terms
|
Cookie Preferences
|

© 2024, Amazon Web Services, Inc. or its affiliates. All rights reserved.

